

MVP Examples

Purpose

Show how successful products started with a *sharply constrained* first version to validate the riskiest assumption fast. Use these examples to shape your own MVP's scope, metrics, and decision gates.

What is an MVP?

A **minimum viable product** is the *smallest artifact or experience* that tests a **single, riskiest assumption** with real users—measured by behavior, not opinions. It can be a video, a landing page, a concierge service, or a single-feature app. The MVP is not a “tiny product”; it’s a **focused experiment**.

Five classic MVPs (patterns to copy)

1) Dropbox — Demo Video as MVP

- **Core bet:** People want seamless file sync across devices; the magic is credibility and clarity.
- **MVP artifact:** A short demo video “showing” the experience (before the product fully existed).
- **What it proved:** *Value risk*—a large audience signaled intent (waitlist signups, email replies) when they understood the promise.
- **Key metric:** Qualified signups and referral velocity immediately after viewing.

Pattern to copy: If your product’s value is hard to grasp until seen, use a tight **Storyboard + screencast** to validate *interest* before building infrastructure.

2) Airbnb — Simple Website & Manual Ops

- **Core bet:** Travelers will pay to stay in a stranger’s home; hosts will list if demand exists.

- **MVP artifact:** A basic website listing a few local spaces; founders did photography and messaging by hand.
- **What it proved:** *Behavioral demand* on both sides of a tiny marketplace (hosts list, guests book).
- **Key metric:** Completed bookings per week in one city; % hosts willing to relist.

Pattern to copy: Start with **one city + one use case**, run the “marketplace glue” **manually** (photos, vetting) to test liquidity before scaling.

3) Instagram — Single Feature App

- **Core bet:** Instant, beautiful mobile photo sharing would drive daily use.
- **MVP artifact:** Photo capture + filters + feed; everything else cut.
- **What it proved:** *Engagement loop*—creation → sharing → likes → return.
- **Key metric:** Day-1/Day-7 retention and photos/user/day.

Pattern to copy: Pick **one atomic action** that users repeat often. Ship *only* what makes that action delightful end-to-end.

4) Twitter — Status Updates (Microblog)

- **Core bet:** Public, short status updates create ambient community.
- **MVP artifact:** SMS/web posting of 140-character messages; minimal discovery.
- **What it proved:** *Network behavior*—people post and check frequently if friction is nearly zero.
- **Key metric:** Posts/day and % of users posting weekly.

Pattern to copy: Reduce the **unit of creation** to the smallest possible action; bias for **real-time** feedback that forms habit.

5) Facebook — Restricted College Directory

- **Core bet:** A real-name, authenticated directory is valuable inside a single campus.
- **MVP artifact:** Basic profiles, friend connections, and a .edu email gate (Harvard only).
- **What it proved:** *Density matters*—utility skyrockets with a tight, trustable graph.

- **Key metric:** Penetration rate on one campus (e.g., % of student body signed up) and daily actives.

Pattern to copy: Gate the audience to a place with natural trust and repeat contact; nail saturation in one micro-market before expanding.

Cross-cutting patterns (what all five did right)

- **One core feature/flow:** demo → booking → photo share → post → profile graph.
- **Tight constraints:** one device, one city, one community, one action.
- **Manual where it's expensive to automate:** “Wizard-of-Oz” operations de-risk before code.
- **Proof = behavior:** signups, bookings, posts, retention—not survey compliments.
- **Narrative clarity:** a crisp promise users “get” in seconds.
- **Expansion path built-in:** once core loop works in a wedge, expand adjacently (next city, campus, feature).

MVP types (choose based on your riskiest assumption)

- **Value risk (Do they care?)** → *Demo video, landing page (fake door), waitlist, paper prototype.*
- **Feasibility risk (Can it work?)** → *Technical spike, concierge/wizard-of-oz with narrow scope.*
- **Growth risk (Will it spread?)** → *Referral loop test, invite-only cohort, seeded community.*
- **Monetization risk (Will they pay?)** → *Time-boxed paid pilot, priced pre-order, WTP survey tied to behavior.*
- **Marketplace liquidity (Will both sides show up?)** → *Single-city manual brokerage, hand-curated supply/demand.*

Designing your MVP (copy-ready checklist)

1. Define the core bet (one sentence)

"If we [enable one action/outcome] for [specific ICP] in [narrow context], then they will [measurable behavior] within [timeframe]."

2. Pick an MVP type that *directly* tests that bet (from the list above).
 3. Constrain scope (platform, audience, locale, content).
 4. Choose success metrics & gates (examples below).
 5. Plan manual steps you'll replace later (what you'll do by hand).
 6. Document ethics & consent (what you collect, how you use it).
 7. Ship in days, not months; instrument basic analytics.
-

Success metrics & decision gates (examples you can adapt)

- Interest/intent (video or landing page):
 - Gate: $\geq 15\text{--}20\%$ of qualified visitors join waitlist; $\geq 30\%$ of waitlisted respond to follow-up.
- Engagement (single-feature app):
 - Gate: $\geq 40\%$ Day-1 retention and $\geq 20\%$ Day-7; median **2+** core actions/day/user.
- Conversion (marketplace city #1):
 - Gate: $\geq 10\%$ of visitors attempt to book; $\geq 60\%$ host relist rate after first booking cycle.
- Monetization (paid pilot):
 - Gate: $\geq 3\text{--}5$ pilot commitments at target ARPU band within **4–6 weeks**.

If you miss gates, either **tighten the wedge, reduce friction** on the core action, or **re-articulate the promise**—then retest.

Do / Don't

Do

- Write the **core bet** first; let it kill nice-to-have features.
- **Over-constrain** (one campus/city/device) to reach signal faster.

- Replace code with **human effort** to learn cheaply.
- Measure **one success metric** tied to the bet (e.g., bookings, posts, retention).
- Plan your **next adjacent move** *only after* the core loop works.

Don't

- Stack multiple risky bets into one build.
 - Treat MVP like a v1 product spec.
 - Chase vanity metrics (views, likes) without downstream behavior.
 - Expand geography/segments before saturation in your wedge.
 - Skip consent or mishandle user data.
-

Common pitfalls & fixes

- **Too many features:** *Fix:* Ship only the **core loop**; defer accounts, settings, dashboards.
- **Vague promise:** *Fix:* Rewrite landing/video headline until a stranger “gets it” in 5 seconds.
- **No wedge:** *Fix:* Pick **one** micro-market (school, city, role).
- **Automating too soon:** *Fix:* Do ops manually for the first 20–50 users.
- **Measuring the wrong thing:** *Fix:* Tie metrics to the **bet** (e.g., bookings, not pageviews).